

# Semi-parametric dynamic contextual pricing

*Virag Shah* — Jose Blanchet — Ramesh Johari  
Stanford University  
virag@stanford.edu

September 27, 2019

# Dynamic pricing

- ▶ Several platforms have access to data describing history of different users.
- ▶ Platforms can leverage this information to price different products and optimize revenue.
- ▶ This requires learning online the mapping from user context to optimal price, efficiently.

# Basic Framework

- ▶ Discrete times  $1, 2, \dots, n$ , one user arrives per time step
- ▶ Each user is shown one product, which is ex-ante fixed
- ▶ Let  $V_t$  be the value  $t^{\text{th}}$  user assigns to the product.
- ▶ Let  $p_t$  be the price set by the platform.
- ▶ The user buys the product if  $p_t \leq V_t$ .
- ▶ Platform does not know or observe  $V_t$ , but has access to covariates  $X_t \in \mathbb{R}^d$  which may describe user's history and product's type
- ▶ Goal: set prices  $p_1, \dots, p_n$  so as to maximize  $\sum_{t=1}^n p_t \mathbf{1}\{p_t \leq V_t\}$ .

# Predicting $V_t$ : The Data

- ▶ Input:  $\{X_i, p_i\}_{i=1}^{t-1}$ .  
 $X_t$  : covariate.  $p_t$ : price

- ▶ Output:  $\{Y_i\}_{i=1}^{t-1}$ , where

$$Y_i = \begin{cases} 1 & \text{if } V_i - p_i \geq 0 \\ 0 & \text{otherwise} \end{cases} .$$

In other words,  $Y_i$  captures whether  $i^{\text{th}}$  was success or failure.

# Predicting $V_t$ : The Model

- ▶ Suppose,

$$V_i = \theta_0^\top X_i + Z_i,$$

$\theta_0$ : unknown parameters,  $Z_i$ : unknown demand-shock/noise.

- ▶ If we assume  $Z_1, \dots, Z_{t-1}$  are i.i.d. **Gaussian**, then we obtain **Probit model** with input  $\{X_i, p_i\}_{i=1}^t$ .
- ▶ Similarly, we obtain Logistic model for i.i.d. Logistic distributed demand-shocks.

# Non parametric noise

- ▶ In this talk, we will not make any Gaussian or Logistic noise assumptions.

# Non parametric noise

- ▶ In this talk, we will not make any Gaussian or Logistic noise assumptions.
- ▶ Even if we knew  $\theta_0$  exactly, the actual structure of distribution of demand-shocks is important to set optimal markups.

# Non parametric noise

- ▶ In this talk, we will not make any Gaussian or Logistic noise assumptions.
- ▶ Even if we knew  $\theta_0$  exactly, the actual structure of distribution of demand-shocks is important to set optimal markups.
- ▶ Can we learn the actual distribution of shocks along with  $\theta_0$ ?



# Predicting $V_t$ : Our Semi-parametric Model

- ▶ Let

$$\ln V_i = \theta_0^\top X_i + Z_i,$$

$\theta_0$ : unknown parameters,  $Z_i$ : unknown demand-shock/noise.

- ▶ Demand-shocks are i.i.d. with arbitrary (non-parametric) distribution.
- ▶ Input:  $\{X_i, p_i\}_{i=1}^{t-1}$ , Output:  $\{Y_i\}_{i=1}^{t-1}$  where  $Y_i = \mathbf{1} \{V_i - p_i \geq 0\}$ .

# Semi-parametric Model: Curse of Dimensionality

- ▶ Usually, such semi-parametric models with binary feedback are hard to learn.
- ▶ That is one of the reasons why previous works on dynamic contextual pricing only consider probit/logistic/generalized-linear models.

# Semi-parametric Model: Curse of Dimensionality

- ▶ Usually, such semi-parametric models with binary feedback are hard to learn.
- ▶ That is one of the reasons why previous works on dynamic contextual pricing only consider probit/logistic/generalized-linear models.
- ▶ Key leverage we find and use: Platform controls prices, and it is possible to focus learning in the regions which are revenue maximizing.

# Price Experimentation

Conflicting goals:

- ▶ Observe  $X_t$ . Set price  $p_t$  to better learn  $\theta_0$  and distribution of  $Z$ . (exploration)
- ▶ Observe  $X_t$ . Set price  $p_t$  to maximize revenue. (exploitation)

# Price Experimentation

Conflicting goals:

- ▶ Observe  $X_t$ . Set price  $p_t$  to better learn  $\theta_0$  and distribution of  $Z$ . (exploration)
- ▶ Observe  $X_t$ . Set price  $p_t$  to maximize revenue. (exploitation)

We view this as contextual bandit problem with prices as arms to maximize revenue.

# Price Experimentation

Conflicting goals:

- ▶ Observe  $X_t$ . Set price  $p_t$  to better learn  $\theta_0$  and distribution of  $Z$ . (exploration)
- ▶ Observe  $X_t$ . Set price  $p_t$  to maximize revenue. (exploitation)

We view this as contextual bandit problem with prices as arms to maximize revenue.

As we learn more, we can zero in the further exploration near profit maximizing regions.

# The Oracle

Let us first consider a simpler problem.

- ▶ Consider an Oracle which actually knows  $\theta_0$  and the distribution of  $Z$ .
- ▶ Now, let

$$F(z) = z\mathbb{P}(Z \geq \ln z), \text{ and } z^* = \arg \sup_z F(z).$$

## Proposition

*The following pricing policy maximizes revenue for the Oracle: At each time  $t$  set price  $p_t^*$  such that*

$$\ln p_t^* = \theta_0^\top X_t + \ln z^*.$$

# Regret

- ▶ Platform's revenue:  $\Gamma_n = \sum_{t=1}^n p_t \mathbf{1}\{p_t \leq V_t\}$ .
- ▶ Oracle's expected revenue:  $\Gamma_n^* = \sum_{t=1}^n p_t^* \mathbf{1}\{p_t^* \leq V_t\}$ .
- ▶ Expected regret:  $\mathbb{E}[R_n] = \mathbb{E}[\Gamma_n^*] - \mathbb{E}[\Gamma_n]$ .
- ▶ Maximizing  $\mathbb{E}[\Gamma_n]$  is equivalent to minimizing  $\mathbb{E}[R_n]$ .
- ▶ We will focus on minimizing  $\mathbb{E}[R_n]$ , asymptotically in  $n$ .



# Designing Optimal Algorithm: Key Ideas

- ▶ Recall, revenue maximizing policy for Oracle:  $\ln p_t^* = \theta_0^\top X_t + \ln z^*$ .
- ▶ For each  $z$  and  $\theta$ , think of  $(z, \theta)$  as an arm (i.e. a potential option). Pulling arm  $(z, \theta)$  is equivalent to setting price  $p_t$  such that  $\ln p_t = \theta^\top X_t + \ln z$ .
- ▶  $(z, \theta) \in \mathbb{R}^{d+1}$ : Curse of dimensionality?
- ▶ Important observation: Given  $X_t$ , for each choice of price  $p_t$  we *simultaneously* obtain information about the expected revenue for a *range* of pairs  $(z, \theta)$ .

# DEEP-C Pricing Algorithm: Summary

DEEP-C: Dynamic Experimentation and Elimination of Prices - with Covariates.

- ▶ Maintain a set  $A(t)$  of 'active arms'  $(z, \theta)$  at each time.
- ▶ At time  $t$ , observe  $X_t$  and compute the set of active prices:

$$P(t) = \{p_t : \exists (z, \theta) \in A(t) \text{ s.t. } \ln p_t = \theta^\top X_t + \ln z\}.$$

- ▶ Choose price  $p_t$  at random from  $P(t)$ .
- ▶ Observe the revenue obtained. Eliminate  $(z, \theta)$ 's from  $A(t)$  for which there is enough information about sub-optimality.

# The main result

Under some smoothness, compactness, independence, etc. assumptions, the following holds.

## Theorem

*The expected regret of DEEP-C algorithm satisfies the following: there exists a constant  $c$  such that*

$$\mathbb{E}[R_n] = O(d^c \sqrt{n}) .$$

# Conclusions

- ▶ To learn via price experimentation, we do not need to make parametric (probit/logistic type) assumptions.
- ▶ We have a provably efficient algorithm which works under a 'very general' setting.